BRILWORKS · EXECUTIVE AUDIT REPORT

SEO, Technical Architecture & Generative Engine Optimization Audit

Evidence-based review of brilworks.com: 16 months of Search Console data, a full crawl of 726 sitemap URLs, the Next.js codebase, and AI-search readiness — with a phased execution roadmap.

58/100
Overall SEO health

−67%
Organic clicks, first vs last 90 days

2.1%
Of clicks reach service, hire or industry pages

521
Blog posts served uncached on every request

https://www.brilworks.com/ Audit date: 11 November 2026 For development, content and executive teams

0. Executive Summary

Overall health: 58 / 100 — Needs Improvement. Organic search is in sustained decline and is not producing commercial demand. Search Console shows clicks down **67%** and impressions down **66%** between the first and last 90 days of the 16-month export. Of 31,953 clicks, service, hire and industry pages together earned **694 (2.1%)**. Queries about Snowflake, Databricks and data engineering — the pillar the company now leads with — earned **7 clicks** among the top 1,000 queries.

The content base is large and the new positioning is clear on the homepage, but the technical layer is working against it. The pages that carry the new strategy are the least discoverable on the site, and the 521 blog posts that earn most visibility are served in the most fragile way possible.

Search performance at a glance (GSC, 10 Jul 2025 – 9 Nov 2026)

Metric	Value
Total clicks / impressions / CTR	31,953 / 11.58 M / 0.28%
Clicks per day, first 90 days → last 90 days	107 → 35 (−67%)
Clicks per day, October 2025 → October 2026	99 → 26 (−74%)
Share of clicks from India	66% (21,200)
United States	48% of impressions, 7% of clicks, CTR 0.04%, average position 31
Clicks to homepage + careers page	30% of all clicks
Clicks to blog	62%
Clicks to service, hire and industry pages	2.1%
Branded share of clicks among the top 1,000 queries	56% (7,357 of 13,131)

The six findings that matter most

#	Finding	Evidence	Why it matters
0	**Traffic is falling, and what remains is brand, careers and developer tutorials from the wrong market**	GSC figures above. Top non-brand queries: "lambda expression in java", "stream api in java", "meme gen ai". The 19 hire pages hold 152,000 impressions at positions 45–68 and earned 153 clicks	The site ranks for what students search, not what buyers search. US buyers see Brilworks on page 3 or later
1	**Every blog post is rendered on demand and never cached**	521 of 521 posts return `cache-control: private, no-store` and `x-vercel-cache: MISS` ; every other template is prerendered	Each Googlebot and AI-crawler hit calls Storyblok live. One CMS slowdown or rate limit takes 72% of the indexable site down for crawlers at once
2	**Any URL under `/blog/` or `/ebooks/` returns HTTP 200**	`/blog/zzz-this-does-not-exist/` → 200 with a `noindex` tag; 13 real internal links point at such dead posts	Unlimited soft-404 space. Crawl budget is spent on pages that do not exist and deleted posts never drop out cleanly
3	**The new strategic pages are orphaned or missing from the sitemap**	`/databricks-partner/` , `/snowflake-partner/` and 3 of 4 case studies have zero internal links; 12 Data & AI service pages, 4 comparison pages and the case-study hub are absent from `sitemap.xml`	The pages meant to win "Databricks consulting partner"-type queries receive almost no internal authority

#	Finding	Evidence	Why it matters
4	**The blog index contains no links to blog posts in its HTML**	`/blog/` server HTML has 0 post links; listing is fetched client-side. 84 posts have no inbound internal link at all	Posts are discoverable only through the sitemap. AI crawlers, which do not run JavaScript, cannot browse the blog
5	**The site describes Brilworks three different ways**	Homepage: "Data & AI Engineering Company". 133 pages carry `og:site_name` "AWS Consulting Partner | Gen AI | Product Engineering". `WebSite` schema: "mobile app and software development company"	Search engines and LLMs build an entity from repeated, consistent statements. The site currently contradicts its own `llms.txt`

Scorecard

Category	Weight	Score	Basis
Technical SEO	25%	55	Full crawl + code review (confirmed)
Content quality	20%	68	Crawl metrics + sampled reading
On-page SEO	15%	62	Full crawl (confirmed)
Schema / structured data	15%	45	Full crawl + code review (confirmed)
Performance (CWV)	10%	60	Lab only — field data not available
Image optimisation	10%	55	Code + crawl
AI search readiness	5%	60	Crawl + third-party search

Scores are a judgement aid, not a measurement. The findings below are what to act on.

What this audit could and could not verify

Source	Status
Live crawl of all 726 sitemap URLs plus 153 linked or unlisted URLs	Done
Codebase (Next.js 16 App Router, Storyblok, Vercel behind Cloudflare)	Done
Lab performance, throttled mobile Chromium, 7 templates	Done — first-party resources only
Google Search Console — Performance export	Done. Web search, 10 Jul 2025 – 9 Nov 2026: daily chart, top 1,000 queries, top 1,000 pages, countries, devices, search appearance
Google Search Console — Indexing, Crawl stats, query-by-page	**Not available.** Not in the export; live browser access to Search Console was not possible from the audit environment
PageSpeed Insights / CrUX field data	Not available — the shared API quota was exhausted
Third-party entity profiles	Search-result level only; profiles were not opened individually

GSC exports are capped at 1,000 rows: the query file covers 41% of all clicks, the page file covers nearly all. Index-coverage statements come from the crawl, not from GSC, and are marked as expected rather than observed.

Part 1 — Technical SEO, Search Console & Crawl Budget

1.1 Google Search Console

Source: Performance export for `sc-domain:brilworks.com`, Web search, 10 July 2025 – 9 November 2026.

Trend: a 16-month decline

Month	Clicks	Impressions	CTR	Avg. position	Clicks / day
Aug 2025	2,995	1,212,096	0.25%	38.2	97
Oct 2025	3,056	724,726	0.42%	18.8	99
Dec 2025	2,686	1,010,698	0.27%	21.2	87
Feb 2026	1,917	761,140	0.25%	18.7	68
Apr 2026	987	418,018	0.24%	25.2	33
May 2026	813	541,674	0.15%	28.0	26
Jul 2026	1,743	517,009	0.34%	29.5	56
Aug 2026	1,579	507,525	0.31%	35.9	51
Sep 2026	1,141	337,985	0.34%	25.4	38
Oct 2026	799	321,429	0.25%	27.2	26

Clicks fell steadily from December 2025, dropped sharply in April–May 2026, recovered partly in July–August, and are back at the low. The export alone cannot prove the cause. Three candidates line up with what the crawl and code show, and all three are addressed in the roadmap:

- **Repositioning without redirecting authority.** The homepage and navigation moved to Data & AI in 2026 while the pages that held rankings (Java, React, Node, mobile) lost prominence and the new pages launched without sitemap entries or internal links.
- **Informational queries losing clicks to AI answers.** The traffic base is definitional developer content ("what is a vector database", "lambda expression in java"), the query type most exposed to AI Overviews. CTR on page-one positions is far below normal (table below).
- **Technical fragility of the blog.** Uncached rendering, soft 404s and an uncrawlable blog index all date from the Next.js migration period.

Where the clicks come from

Segment	Clicks	Share	Impressions	CTR
Blog (790 URLs)	20,250	62.1%	11.81 M	0.17%
Homepage	6,188	19.0%	46,774	13.2%
`/career/`	3,727	11.4%	25,745	14.5%
AI solutions + AI agent demos	1,015	3.1%	82,447	1.2%
Service pages	**508**	**1.6%**	114,725	0.44%

Segment	Clicks	Share	Impressions	CTR
Free tools	284	0.9%	40,123	0.71%
Hire pages (19)	**153**	**0.5%**	151,972	0.10%
Use-case pages	105	0.3%	154,260	0.07%
Portfolio	59	0.2%	9,409	0.63%
Industry pages (8)	**33**	**0.1%**	23,240	0.14%
Data-engineering tools	11	0.0%	1,178	0.93%

The ten most-clicked pages account for 45% of clicks; 723 of the 1,000 pages earned fewer than ten clicks in 16 months, and 177 pages have over 1,000 impressions with zero clicks.

Who is searching

Country	Clicks	Impressions	CTR	Avg. position
India	21,200	2,272,406	0.93%	12.5
United States	2,184	5,616,046	0.04%	31.4
United Kingdom	509	524,765	0.10%	37.4
Canada	387	231,780	0.17%	25.8
Australia	239	148,863	0.16%	39.6

Two-thirds of clicks are from India, much of it brand and careers ("brilworks technology private limited", "brilworks careers"). In the United States — the main buyer market — the site averages position 31 and converts 0.04% of impressions.

What people search for

Query group (top 1,000 queries)	Clicks	Impressions
Brand ("brilworks" and variants)	7,357	18,771
Java	1,785	386,936
AI / generative AI / agents	1,091	336,454
React / React Native	839	409,153
Node.js	435	54,841
AWS / cloud	198	157,583
Mobile / app	130	111,165
Commercial-intent queries (hire, company, services, consulting, cost) — 135 queries	**470**	228,446
Snowflake / Databricks / data engineering	**7**	2,557

Brand searches use four names — "brilworks", "brilworks software", "brilworks technology private limited", "brilworks technology pvt ltd" — the same naming inconsistency found in third-party profiles (section 3.1).

Commercial queries: visible but far from page one

Query	Impressions	Clicks	Position
aws consulting	20,182	1	49.2
aws consulting services	14,567	1	47.5
java development services	13,024	1	52.5
aws consultant	12,875	4	26.3
hire aws developers	7,997	4	43.9
hire java developers	6,247	2	62.5
nodejs development company	6,081	1	40.9
hire nodejs developers	6,020	4	45.2
hire reactjs developers	4,439	1	52.7
react native app development company	3,550	2	54.3
outsource java development	3,282	2	25.6

Google associates Brilworks with these queries but ranks it on pages 3–7. This is the pattern expected when commercial pages have thin internal linking and weak external authority, which is what sections 1.4 and 3.1 found.

Commercial pages closest to page one

These pages already rank between positions 8 and 25 and are the fastest commercial wins.

Page	Impressions	Clicks	Position	Action
`/ai-ml-development-services/`	9,823	20	23.3	Query-led H1, add table and case study, link from the 60 AI posts
`/industry/healthcare-software-development/`	3,583	9	16.6	Link from the 21 healthcare posts; add `Service` schema
`/services/ai-agent-development/`	2,786	5	15.2	Duplicate of the main AI agent page — 301 and consolidate
`/application-development-services/`	2,777	9	9.2	Title and description rewrite for CTR
`/product-engineering-development-services/`	1,584	3	8.7	Title rewrite; make it the silo hub
`/embedded-software-development-services/`	1,301	1	10.2	Add missing canonical and schema; link from homepage
`/industry/manufacturing-software-development-services/`	1,301	0	8.6	Title and description rewrite
`/mobile-app-development-services/`	1,048	4	9.8	Query-led H1
`/hire-white-label-developer/`	9,326	55	35.3	Not in sitemap; add it
`/use-case/aws-in-healthcare/`	31,436	56	21.2	Add canonical; link to healthcare and AWS service pages

The eight `/use-case/*` pages hold 154,000 impressions between positions 21 and 75 with no canonical tag and at most three contextual inbound links each.

Quick wins: high impressions, low CTR, already on or near page one

Page	Impressions	Clicks	CTR	Position	What to do
`/blog/what-is-vector-database-types-uses/`	416,433	153	0.04%	14.4	Answer-first definition, comparison table of vector DBs, refresh; target top 5
`/blog/apple-human-interface-guidelines/`	415,706	189	0.05%	8.9	Navigational intent for Apple's own page; low value — do not invest
`/blog/best-artificial-intelligence-platforms/`	338,544	598	0.18%	21.6	Update for 2026, add a ranked table; "ai platforms" sits at 16.5
`/blog/react-vs-vue/`	231,841	22	0.01%	4.0	One query shows position 1.1 with 210,174 impressions and 2 clicks — an impression pattern typical of a SERP feature rather than a normal listing. Inspect the live SERP before investing
`/blog/path-to-become-aws-partner/`	205,188	403	0.20%	15.8	Split title toward "AWS partner tiers and requirements" (positions 7–9); link to `/aws-consulting-services/`
`/blog/lambda-expression-java/`	168,763	842	0.50%	8.7	Tighten title; code example in the first screen; FAQ already opted in
`/blog/whats-new-in-spring-boot-3-for-java-developers-in-2023/`	116,277	109	0.09%	11.0	Year-stamped 2023 — retitle to the current Spring Boot version and refresh
`/blog/iot-in-mobile-app-development/`	107,535	27	0.03%	13.9	Rewrite title and intro; link to the IoT service pages
`/blog/nestjs-vs-nextjs/`	101,529	345	0.34%	8.2	Comparison table at top; verdict sentence under H1
`/blog/javascript-web-frameworks-comparison/`	97,100	117	0.12%	13.5	Make it the frameworks hub (section 2.3)
`/blog/react-native-vs-kotlin/`	92,652	77	0.08%	9.6	Ranks for "what is kotlin" (69,720 impressions, 1 click): intent mismatch. Add a definitional section or accept
`/blog/Java-21-features/`	66,606	219	0.33%	8.6	Mixed-case URL; update with Java 25 cross-link

Expect modest gains from the developer-tutorial pages. They bring students and engineers, mostly from India, and rarely buyers. Their real value is as internal-link donors: these twelve pages hold 2.4 million impressions, yet six contain no link to any service page, one links to a page that returns 404 (`/ios-app-development/`), and one links to the non-`www` domain through redirects.

Technical issues confirmed by GSC

Issue	GSC evidence
Dead post still served as 200	`/blog/overcomming-common-generative-ai-challenges/` — 14,898 impressions, 18 clicks, now a soft 404. Redirect to `/blog/challenges-of-generative-ai/`
Dead post still served as 200	`/blog/what-is-databricks-a-complete-guide-for-modern-data-teams/` — 1,576 impressions. Redirect to `/blog/what-is-databricks/`
Redirected URLs still shown in results	10 URLs, 15,323 impressions, e.g. `/blog/how-to-hire-node-js-developers/` (9,916). Internal links to the old URLs keep them alive
Indexable pages outside the sitemap earning impressions	37 URLs, 37,842 impressions: `/product/white-label-fitness-app/` , `/hire-white-label-developer/` , author pages, `/login/` (1,169 impressions)
`/login/` indexed	Add `noindex`
App subdomains under the same property	15+ hosts such as `eyerest.` , `roomplanner.` , `coursesphere.brilworks.com` appear in results. Confirm each is intended to be public; they dilute the domain's topical focus
"Review snippet" search appearance	438,533 impressions at average position 34.7. No review markup exists in the current `www` JSON-LD or code, so the source is unidentified (possibly CMS-embedded markup or a subdomain). Find it; self-serving review markup risks a manual action
Mobile share	Mobile is 21% of impressions and 31% of clicks — a developer, desktop-heavy audience

Still to pull from Search Console

The Performance export does not include indexing or crawl data. These three exports complete the picture:

Report	Purpose
Pages (indexing) → every reason, exported	Confirm the expected reasons below
Settings → Crawl stats → by response and by purpose	Share of 404 and 308 responses; response time on `/blog/`
Performance → Queries filtered to each top commercial query → Pages tab	Confirm the cannibalisation clusters in section 2.3

What the indexing report is expected to show

GSC reason	Cause found in the crawl
Soft 404	Dead `/blog/*` and `/ebooks/*` URLs returning 200
Excluded by 'noindex'	The same dead URLs, plus 3 `noindex` URLs listed in the sitemap
Page with redirect	34 internal links without a trailing slash, the footer `//` link, links to `brilworks.com` without `www`
Not found (404)	10 internal links to missing pages, 1 sitemap URL, 4 video-sitemap entries
Duplicate without user-selected canonical	27 sitemap pages with a missing or wrong canonical

GSC reason	Cause found in the crawl
Crawled – currently not indexed	31 `/ai-agents/*` pages with two words of server HTML; thin `/tools/` and `/gist/` pages
Discovered – currently not indexed	Pages reachable only through the sitemap, including 84 orphan posts

1.2 Rendering strategy and Core Web Vitals

How each template is served today

Template	URLs	Rendering	Cache	Verdict
Homepage, services, hire, industry, tools	~150	Static prerender	CDN HIT	Good
Portfolio, ebooks, use cases	~33	ISR	HIT / STALE	Good
Blog posts	**521**	**Dynamic SSR on every request**	**`no-store`, MISS**	**Fix first**
Blog index	1	Static shell, list fetched in browser	HIT	Content invisible to non-JS crawlers
`/ai-agents/*`	31	Static shell, app runs in browser	HIT	2 words of HTML
`robots.txt`	1	Dynamic (reads host header)	—	Acceptable

Why the blog is dynamic

`src/app/blog/[slug]/page.js` defines `generateStaticParams`, but its data fetch uses `next: { revalidate: 0 }`. A zero revalidate opts the whole route out of static generation, so the static params are never used. Two side effects follow:

- `generateStaticParams` calls `getblog()`, which requests only the first 100 posts, so 421 posts would not be prebuilt even if static generation worked.
- Because the response has already started streaming when `notFound()` runs, the status stays 200 and only a `noindex` tag is added. This is the source of the soft 404s.

Fix: move the blog to ISR with on-demand revalidation

```
// src/app/blog/[slug]/page.js
export const revalidate = 86400;      // safety net: rebuild at most daily
export const dynamicParams = true;    // new posts render on first request, then cache

export async function generateStaticParams() {
  const posts = await getBlogForSitemap();          // already paginates all pages
  return posts.map((p) => ({ slug: p.loc.split("/blog/")[1].replace(/\/$/, "") }));
}

export async function fetchData(slug) {
  const res = await fetch(storyUrl, {
    next: { revalidate: 86400, tags: [`blog:${slug}`] },   // was: revalidate: 0
  });
  if (res.status === 404) return null;                       // lets notFound() return a real 404
  ...
}
```

The Storyblok webhook at `/api/revalidate` already calls `revalidatePath("/" + full_slug)`, so published edits will still appear within seconds. After deploying, confirm with:

```
curl -sI https://www.brilworks.com/blog/node-js-vs-golang/ | grep -i "x-vercel-cache\|x-nextjs-prerender"
```

```
curl -s -o /dev/null -w "%{http_code}\n" https://www.brilworks.com/blog/zzz-this-does-not-exist/
```

Expected: `HIT` with `x-nextjs-prerender: 1`, and `404`. Apply the same change to `src/app/ebooks/[slug]/page.js`.

Lab performance (mobile emulation, 4× CPU slowdown, slow 4G, first-party only)

Template	FCP	LCP	LCP element	CLS	Blocking time	Scripts
Homepage	2.0 s	2.0 s	Hero paragraph	0.01	**1,340 ms**	51
`/data-engineering-services/`	1.5 s	1.5 s	Hero paragraph	0.05	860 ms	50
Blog post	1.5 s	1.5 s	H1	**0.38**	740 ms	56
`/hire-reactjs-developer/`	1.7 s	**5.4 s**	Hero image	0.01	890 ms	63
Industry page	1.2 s	1.2 s	Hero paragraph	0.01	900 ms	61
Portfolio case study	1.6 s	1.6 s	Hero paragraph	0.01	700 ms	52
Tool page	1.6 s	1.6 s	Hero paragraph	0.01	780 ms	60

These are lab numbers with third-party scripts blocked, so real blocking time is higher. The blog CLS figure was measured with CMS images blocked and needs confirming in the field. Treat this table as direction, and confirm with CrUX in Search Console → Core Web Vitals.

Recommendations by metric

INP / blocking time (the main risk on every template)

- 50–63 script files and 33–39 background fetches load per page. Most fetches are Next.js link prefetches for the 64 sitewide navigation and footer links. Add `prefetch={false}` to footer and mega-menu links; only 3 links in the codebase do this today.
- Five third-party tags load on every page (GTM, PostHog, Clarity, Factors, Clearbit, Vector, Cloudflare Insights, plus a client-side IP lookup). Audit which are still used. Load survivors through GTM with consent gating.
- Homepage uses `framer-motion`, Swiper and Splide together. Standardise on one carousel library and replace scroll-reveal animation with CSS.
- The header fetches Storyblok data in the browser on every page. Move it to the server component.

LCP

- Hire pages: the hero image is the LCP element and is discovered late (5.4 s). Add `priority` and `fetchPriority="high"` to the hero `<Image>` in the hire-page hero component. 41 of 43 images on that page are lazy, including the one that matters.
- `next.config.js` sets `images.unoptimized: true`. No image on the site gets responsive `srcset`, AVIF, or resizing. Either enable the Next.js optimiser, or use Storyblok's image service (`/m/800x0/filters:format(avif)`) through the existing `formatSrcUrl` helper, with explicit `sizes`.

- Service page section images are 125–146 KB each as fixed-size WebP; `/images/banner1.png` on portfolio pages is 190 KB. `public/images` is 59 MB.
- The homepage preloads five partner logos as high-priority images. Preload only the actual LCP resource.

CLS

- Blog posts: give the banner and in-article images explicit `width` / `height` (or `aspect-ratio`), and reserve height for the "Summarize with AI" bar and author block that mount after hydration.
- 10 of 24 homepage images have no dimensions.

1.3 Sitemap and indexation hygiene

What is in `sitemap.xml` today (726 URLs, one flat file)

Section	URLs	Issue
Blog	522	Fine as a set; lastmod comes from Storyblok
Free tools	41	23 are orphans; 36 have two H1s; 2 have wrong canonicals
`/ai-agents/*` demos	32	31 are near-empty HTML
Services, hire, industry, partner	~70	7 have no canonical
Portfolio, case studies	21	3 case studies orphaned
`/x/product-listing/*`	10	1 returns 404; 2 wrong canonicals
Ebooks, use cases, AI solutions	22	14 have no canonical

Good news: there are no tag archives, pagination URLs or parameter URLs in the sitemap, and all 726 use a trailing slash consistently.

Errors in the sitemap

Issue	URLs
Returns 404	`/x/product-listing/service-squard/` (typo; the live page is `service-squad`)
Listed but `noindex`	`/privacy-policy/`, `/terms-and-conditions/`, `/ebooks/saas-founders-guide-2024/`
Canonical points to a different page	`/tools/which-app-should-you-build/` → `mvp-feature-selector-quiz`; `/tools/mobile-app-monetization-strategy-selector/` → `app-rebuild-vs-refactor-calculator`; `/x/product-listing/pet-finder/` → listing page; `/x/product-listing/enclave-connect/` → misspelt `encalve-connect`
No canonical tag (23 pages)	All 8 `/use-case/*`, all 6 `/ai-solutions/*`, `/embedded-software-development-services/`, `/firmware-engineering-services/`, `/edge-ai-development-services/`, `/iot-product-engineering-services/`, `/data-engineering-for-healthtech/`, `/data-engineering-for-fintech/`, `/data-engineering-for-saas/`, `/x/product-listing/`
Stale hard-coded `lastmod`	Static pages are listed in `constants.js` with fixed dates, many still `2025-04-25` although the pages changed in 2026
Mixed-case URL	`/use-case/aws-in-IoT/` (lower-case version 404s)

Indexable pages missing from the sitemap (about 45)

Group	Pages
Data & AI service pages (12)	`/ai-strategy-advisory-services/`, `/ai-readiness-data-foundation-services/`, `/ai-engineering-aiops-services/`, `/ai-governance-services/`, `/enterprise-data-strategy/`, `/cloud-data-engineering/`, `/data-pipeline-development/`, `/data-modernization/`, `/data-governance-consulting/`, `/business-intelligence-development/`, `/business-intelligence-services/`, `/iot-embedded-development-services/`
Comparison pages (4)	`/compare/brilworks-vs-toptal/`, `/compare/brilworks-vs-upwork-agencies/`, `/compare/build-vs-buy-software/`, `/compare/in-house-vs-agency-development/`
Hubs	`/case-studies/`, `/use-case/`, `/hire-engineers/`, `/portfolio/ai-agents/`, `/ai-tools/`
Work	`/portfolio/claude-code-coding-speed-lift/`, 12 `/x/product-listing/*` pages
Other	3 ebooks, `/hire-white-label-developer/`, `/erp-next-development-services/`, `/ai-tools-audit/`, author pages

The cause is structural: static pages are added to the sitemap by hand in `sitemapData` inside `constants.js`. Every new page is a chance to forget.

Other indexation leaks

- **Three URL patterns for the same author:** `/blog/author/<name>/`, `/authors/<name>/` and `/teams/<name>/` all return 200. Keep `/authors/`, 301 the other two.
- **`video-sitemap.xml`** lists `/pre-built-solutions/` four times; that URL returns 404. The real page is `/ai-tools/pre-build-solutions/`.
- **`/gist/*`** (8 hub pages, 140–175 words each) and **`/labs/`**, **`/posters/`** (0 words) are indexable with no canonical. Either build them out or `noindex` them.
- **Legal pages for client apps** (about 30 `*-privacy-policy` pages) are correctly `noindex`. No action.
- **`/?s=`** returns 200 for any value and the `WebSite` schema advertises it as the site search, yet the site has no search and robots.txt disallows the path. Remove the `SearchAction`.

Recommended `robots.txt`

```
User-agent: *
Allow: /
Disallow: /api/
Disallow: /admin/
Disallow: /dashboard/
Disallow: /internal/
Disallow: /login/
Disallow: /forget-password/
Disallow: /verifyemail/
Disallow: /cdn-cgi/
Disallow: /*?s=
Disallow: /*?utm_
Disallow: /*?service=

# AI search and assistant crawlers: allowed
User-agent: GPTBot
User-agent: OAI-SearchBot
User-agent: ChatGPT-User
User-agent: ClaudeBot
User-agent: Claude-SearchBot
User-agent: Claude-User
User-agent: PerplexityBot
```

```
User-agent: Perplexity-User
User-agent: Google-Extended
User-agent: Applebot-Extended
User-agent: CCBot
Allow: /
Disallow: /api/
Disallow: /admin/
Disallow: /dashboard/

Sitemap: https://www.brilworks.com/sitemap.xml
```

Changes from today: removes four WordPress-era rules that match nothing on the current site (`/search/`, `/testimonial/`, `/category/*`, `/email-sms/*`); blocks the private app areas; adds the search-specific AI agents that are currently unnamed; and re-allows `CCBot`. Common Crawl is a primary training source for several LLMs, so blocking it works against the GEO goal — keep the block only if that was a deliberate legal decision.

Do not disallow the legal pages; they need to be crawlable for their `noindex` to be seen.

Sitemap segmentation

Replace the single file with an index, generated from the route tree and CMS rather than a hand-kept list.

File	Contents	Approx. URLs
`sitemap-services.xml`	Services, hire, industry, partner, comparison, pricing	95
`sitemap-work.xml`	Case studies, portfolio, product listings	45
`sitemap-blog.xml`	Blog posts (lastmod from Storyblok)	521
`sitemap-resources.xml`	Ebooks, use cases, AI solutions, author pages	30
`sitemap-tools.xml`	Calculators and assessments that pass a quality bar	30–50
`video-sitemap.xml`	Corrected video entries	7

Segmenting lets Search Console report indexed-versus-submitted per section, which is the fastest way to see whether service pages or tools are the ones being dropped.

```
// src/app/sitemap.js
export async function generateSitemaps() {
  return [{ id: "services" }, { id: "work" }, { id: "blog" }, { id: "resources" }, { id: "tools" }];
}
export default async function sitemap({ id }) {
  // build each list from a route manifest + Storyblok, filter out noindex pages,
  // and take lastModified from git commit date or the page's LAST_UPDATED constant
}
```

Add a CI check that fails the build when a `page.js` without `robots: { index: false }` is absent from the manifest.

1.4 Internal linking and PageRank flow

How authority flows today

Measure	Value
Links repeated on every page (header + footer)	64

Measure	Value
Blog posts with no contextual link to any commercial page	**293 of 521 (56%)**
Sitemap URLs with zero inbound internal links	**127** (84 posts, 23 tools, 20 others)
Sitemap URLs with one or zero contextual inbound links	282
Internal links that redirect (308)	34 targets + 1 sitewide footer link
Internal links that return 404	10 targets

Where blog posts send their authority

Destination	Posts linking to it	Fits current positioning?
`/hire-mobile-app-developer/`	67	No
`/application-development-services/`	35	No
`/ai-ml-development-services/`	35	Partly
`/aws-consulting-services/`	32	Partly
`/generative-ai-development-services/`	26	Yes
`/data-engineering-services/`	10	Yes — core pillar
`/ai-agent-development-services/`	6	Yes — core pillar
`/databricks-partner/` , `/snowflake-partner/`	**0**	Yes — core pillar
Any `/case-studies/*` page	**0**	Yes — proof

The blog still funds the 2024 positioning. The pages that represent the 2026 strategy receive almost nothing from 521 posts.

Leaks to fix

Leak	Detail
Orphaned money pages	`/databricks-partner/` , `/snowflake-partner/` , `/kubernetes-consulting-services/` , `/edge-ai-development-services/` , `/iot-product-engineering-services/` , `/frappe-development-services/` , `/white-label-mobile-app-development/` , `/event-app-development/` , `/hire-adalo-developer/` , `/rapid-digitalization/` , `/services/ai-agent-development/`
Orphaned proof	3 of 4 `/case-studies/*` , 4 ebooks
Footer typo	`/hire-ui-ux-designer//` on 692 pages (`src/app/components/Footer/index.js:198`) — one redirect per page crawled
Footer weight	Footer "Technologies" column links Java, Node, React, React Native, AWS, UI/UX, Blockchain hire pages sitewide; none of the six Data & AI hire pages get footer weight beyond the header
Dead links in content	`/services/` (404, and listed in the sitemap), `/services/claude-code-implementation-and-enablement/` , `/agentic-ai-software-development-services/` , `/cloud-data-engineering-services/` , `/ios-app-development/` , `/aws-architecture` , `/compliance` , `/disaster-recovery` , `/serverless-development` , `/cloud-managed-services`
Leaked internal URLs	Links to `dev.brilworks.com` (7 targets) and to `/BRI/issues/BRI-56` (an internal ticket) inside published content

Leak	Detail
Duplicate service URL	`/services/ai-agent-development/` and `/ai-agent-development-services/` both live and both in the sitemap
Blog breadcrumb	Links to `/blog` (redirects) and shows the category as plain text with no link

Target architecture: five silos, hub and spoke

```
Homepage
 ├─ Data Engineering  (hub: /data-engineering-services/)
 │    spokes: cloud-data-engineering, data-pipeline-development, data-modernization,
 │            data-governance-consulting, business-intelligence-development,
 │            enterprise-data-strategy, snowflake-partner, databricks-partner
 │    hire:   hire-data-engineer, -snowflake-, -databricks-, -analytics-engineer
 │    proof:  4 case studies · tools: 7 data calculators · blog: ~30 posts
 │
 ├─ AI & Agents       (hub: /ai-agent-development-services/)
 │    spokes: generative-ai, ai-ml, ai-strategy-advisory, ai-readiness, aiops, ai-governance
 │    hire:   hire-ml-ai-engineer, hire-mlops-engineer
 │    proof:  portfolio/ai-agents, ai-solutions · blog: ~60 posts
 │
 ├─ Product Engineering (hub: /product-engineering-development-services/)
 │    spokes: saas, web-app, mobile-app, react-native, reactjs, low-code, mvp-in-48-hours
 │    hire:   reactjs, react-native, nodejs, java, mobile, ui-ux
 │    proof:  portfolio · tools: cost calculators · blog: ~250 posts
 │
 ├─ Cloud & DevOps    (hub: /aws-consulting-services/)
 │    spokes: devops-consulting, kubernetes-consulting · hire: hire-aws-developer
 │
 └─ Industries        (hub: new /industries/ index)
      healthcare ⇄ data-engineering-for-healthtech ⇄ healthcare case study ⇄ ~21 posts
      fintech    ⇄ data-engineering-for-fintech    ⇄ fintech case study
      saas       ⇄ data-engineering-for-saas       ⇄ saas case study
```

Linking rules to implement

1. **Every blog post links to exactly one hub and one spoke** in its silo, in the first 300 words, with descriptive anchor text. Implement as a Storyblok field (`primary_service`) rendered as an inline callout, so it can be set in bulk for the 293 posts that have none.
2. **Every spoke links up to its hub, across to two sibling spokes, and down to one case study and one tool.**
3. **Every hub links to all spokes, all related hire pages, and its three best posts.**
4. **Every case study links to the service it proves and the industry page it belongs to**, and each service page shows its case studies. This is the single cheapest fix for the orphaned case studies.
5. **Add visible breadcrumbs with links** on every template, matching the `BreadcrumbList` schema.
6. **Render the blog index on the server** with paginated, crawlable links (`/blog/page/2/`) and linked category pages (`/blog/category/data-engineering/`). This removes the 84 orphans in one change.
7. **Rebalance the footer** to the five hubs plus partner pages. Sitewide links are the strongest internal signal the site sends.
8. **Normalise links at the source:** a lint rule or Storyblok validation that rejects internal hrefs without a trailing slash, without `www`, or pointing at `dev.`.

1.5 Structured data

Coverage today

Template	Pages	Schema present beyond the global three	Gap
Global (every page)	726	`Organization`, `LocalBusiness`, `WebSite`	Three disconnected blocks, no `@id`, conflicting data
Homepage	1	none	No page-level entity
Service pages	42	`WebPage` on 24, `Service` on 2, `FAQPage` on 2	`Service` missing on 40
Unlisted Data & AI pages	12	`Service` + `FAQPage` + `WebPage` on 10	Good — but not in sitemap
Hire pages	19	Two-item `BreadcrumbList` whose last item has no URL	No `Service`, no offer
Industry pages	8	`WebPage`, `BreadcrumbList` on 6	No `Service`
Blog posts	521	`BlogPosting`	No `BreadcrumbList`; `Person` on 4; estimated `wordCount`
Case studies / portfolio	21	`Article` on 3	No client, industry, result or tech data
Tools	41	none	`WebApplication` helper exists but is used once
Data-engineering tools	8	`WebApplication` + `FAQPage` on 7	Good
Ebooks, use cases	16	none	No `Book` / `Article`

Defects in the global schema (`src/app/components/lib/schemaCode.js`)

Defect	Detail
Malformed email	`Organization.contactPoint.email` is `"<sales@brilworks.com>"` including the angle brackets
Conflicting `sameAs`	`Organization` lists 4 profiles; `LocalBusiness` lists 5 different URL forms. Clutch, GoodFirms, Crunchbase, X and GitHub are missing from one or both
Conflicting phone	Schema: `+91 9313644148`. Footer: `+91 9558288364`
Outdated description	`WebSite.description` says "mobile app and software development company"
Fake search	`SearchAction` on `/?s=` — no search exists
SVG logo	Google requires a raster logo (PNG/JPG/WebP, at least 112×112)
`LocalBusiness` on every page	Should appear once, on homepage and contact/about
No `@id` linking	Nothing ties `BlogPosting.publisher` or `Service.provider` to the main organisation node

A note on **FAQPage**: Google has shown FAQ rich results only for government and health authority sites since August 2023. Marking up FAQs on service pages will not produce rich results. It is still harmless and machine-readable when it mirrors visible content exactly, so treat it as low priority, not as a win.

Production-ready JSON-LD

Values in `REPLACE_…` must be filled with verified facts before shipping. Do not publish placeholder or estimated values.

A. Site-wide entity graph — root layout, replaces all three current blocks

```
{
  "@context": "https://schema.org",
  "@graph": [
    {
      "@type": ["Organization", "ProfessionalService"],
      "@id": "https://www.brilworks.com/#organization",
      "name": "Brilworks",
      "legalName": "Brilworks Software LLP",
      "alternateName": "Brilworks Software",
      "url": "https://www.brilworks.com/",
      "logo": {
        "@type": "ImageObject",
        "@id": "https://www.brilworks.com/#logo",
        "url": "https://www.brilworks.com/images/brilworks-logo-512.png",
        "width": 512,
        "height": 512
      },
      "image": "https://www.brilworks.com/images/v2/og-home.jpg",
      "description": "Brilworks is a data and AI engineering firm. It builds data pipelines, warehouses and
governance on Snowflake and Databricks, then the AI agents and analytics that run on that foundation.",
      "foundingDate": "REPLACE_YYYY",
      "founder": [
        { "@id": "https://www.brilworks.com/authors/hitesh-umaletiya/#person" },
        { "@id": "https://www.brilworks.com/authors/vikas-singh/#person" }
      ],
      "numberOfEmployees": { "@type": "QuantitativeValue", "minValue": 50, "maxValue": 249 },
      "address": {
        "@type": "PostalAddress",
        "streetAddress": "503, Fortune Business Hub, Science City Road, Sola",
        "addressLocality": "Ahmedabad",
        "addressRegion": "Gujarat",
        "postalCode": "380060",
        "addressCountry": "IN"
      },
      "geo": { "@type": "GeoCoordinates", "latitude": 23.071, "longitude": 72.5158 },
      "telephone": "REPLACE_ONE_CANONICAL_NUMBER",
      "email": "sales@brilworks.com",
      "contactPoint": [
        { "@type": "ContactPoint", "contactType": "sales", "email": "sales@brilworks.com", "telephone":
"REPLACE_ONE_CANONICAL_NUMBER", "areaServed": "Worldwide", "availableLanguage": "en" },
        { "@type": "ContactPoint", "contactType": "human resources", "email": "hr@brilworks.com" }
      ],
      "areaServed": ["US", "GB", "CA", "AU", "AE", "IN"],
      "knowsAbout": [
        "Data engineering", "Snowflake", "Databricks", "dbt", "Data governance",
        "AI agent development", "Generative AI", "MLOps",
        "Embedded software", "Firmware engineering", "Edge AI"
      ],
      "hasCredential": [
        { "@type": "EducationalOccupationalCredential", "name": "Databricks Partner", "recognizedBy": {
"@type": "Organization", "name": "Databricks", "url": "https://www.databricks.com/" } },
        { "@type": "EducationalOccupationalCredential", "name": "Snowflake Partner", "recognizedBy": {
"@type": "Organization", "name": "Snowflake", "url": "https://www.snowflake.com/" } },
        { "@type": "EducationalOccupationalCredential", "name": "AWS Partner", "recognizedBy": { "@type":
"Organization", "name": "Amazon Web Services", "url": "https://aws.amazon.com/" } }
      ],
      "sameAs": [
        "https://www.linkedin.com/company/brilworks",
        "https://clutch.co/profile/brilworks-software",
        "https://www.goodfirms.co/company/brilworks-software-llp",
        "https://www.crunchbase.com/organization/brilworks",
```

```
        "https://www.g2.com/products/brilworks-software",
        "https://x.com/_Brilworks",
        "https://www.youtube.com/@brilworks",
        "https://www.facebook.com/brilwork",
        "https://www.instagram.com/brilworkssoftware",
        "REPLACE_GITHUB_ORG_URL",
        "REPLACE_WIKIDATA_ENTITY_URL",
        "REPLACE_OFFICIAL_DATABRICKS_PARTNER_DIRECTORY_URL",
        "REPLACE_OFFICIAL_SNOWFLAKE_PARTNER_DIRECTORY_URL"
      ]
    },
    {
      "@type": "WebSite",
      "@id": "https://www.brilworks.com/#website",
      "url": "https://www.brilworks.com/",
      "name": "Brilworks",
      "description": "Data and AI engineering firm: Snowflake and Databricks data engineering, AI agents,
and embedded engineers.",
      "publisher": { "@id": "https://www.brilworks.com/#organization" },
      "inLanguage": "en-US"
    }
  ]
}
```

Remove any `sameAs` line that cannot be filled with a real URL. Use `aggregateRating` only if the reviews are shown on the same page and collected first-hand; a Clutch score copied into schema is self-serving markup and is not eligible for stars.

B. Service page — one per service URL

```
{
  "@context": "https://schema.org",
  "@graph": [
    {
      "@type": "Service",
      "@id": "https://www.brilworks.com/data-engineering-services/#service",
      "name": "Data Engineering Services",
      "serviceType": "Data engineering consulting",
      "description": "Pipelines, warehousing and data modelling on Snowflake and Databricks, including
migration, cost optimisation and AI-ready data foundations.",
      "url": "https://www.brilworks.com/data-engineering-services/",
      "provider": { "@id": "https://www.brilworks.com/#organization" },
      "areaServed": ["US", "GB", "CA", "AU", "AE"],
      "audience": { "@type": "BusinessAudience", "audienceType": "Heads of Data, CTOs and VPs of Engineering
at mid-market and enterprise companies" },
      "hasOfferCatalog": {
        "@type": "OfferCatalog",
        "name": "Data engineering engagements",
        "itemListElement": [
          { "@type": "Offer", "itemOffered": { "@type": "Service", "name": "Fixed-price data platform audit"
} },
          { "@type": "Offer", "itemOffered": { "@type": "Service", "name": "Snowflake or Databricks
migration", "url": "https://www.brilworks.com/data-modernization/" } },
          { "@type": "Offer", "itemOffered": { "@type": "Service", "name": "Data pipeline development",
"url": "https://www.brilworks.com/data-pipeline-development/" } },
          { "@type": "Offer", "itemOffered": { "@type": "Service", "name": "Embedded data engineers", "url":
"https://www.brilworks.com/hire-data-engineer/" } }
        ]
      }
    },
    {
      "@type": "WebPage",
      "@id": "https://www.brilworks.com/data-engineering-services/#webpage",
      "url": "https://www.brilworks.com/data-engineering-services/",
      "name": "Data Engineering Services | Snowflake & Databricks Experts",
      "isPartOf": { "@id": "https://www.brilworks.com/#website" },
```

```
      "about": { "@id": "https://www.brilworks.com/data-engineering-services/#service" },
      "breadcrumb": { "@id": "https://www.brilworks.com/data-engineering-services/#breadcrumb" },
      "dateModified": "2026-04-23"
    },
    {
      "@type": "BreadcrumbList",
      "@id": "https://www.brilworks.com/data-engineering-services/#breadcrumb",
      "itemListElement": [
        { "@type": "ListItem", "position": 1, "name": "Home", "item": "https://www.brilworks.com/" },
        { "@type": "ListItem", "position": 2, "name": "Data & AI Services", "item":
"https://www.brilworks.com/services/" },
        { "@type": "ListItem", "position": 3, "name": "Data Engineering Services", "item":
"https://www.brilworks.com/data-engineering-services/" }
      ]
    }
  ]
}
```

Position 2 assumes a `/services/` hub exists; it currently returns 404 and should be built (section 2.1).

C. Hire page — add an indicative offer

```
{
  "@context": "https://schema.org",
  "@type": "Service",
  "@id": "https://www.brilworks.com/hire-data-engineer/#service",
  "name": "Hire Data Engineers",
  "serviceType": "IT staff augmentation — data engineering",
  "provider": { "@id": "https://www.brilworks.com/#organization" },
  "areaServed": ["US", "GB", "CA", "AU"],
  "url": "https://www.brilworks.com/hire-data-engineer/",
  "offers": {
    "@type": "Offer",
    "priceSpecification": {
      "@type": "UnitPriceSpecification",
      "price": "REPLACE_MONTHLY_STARTING_RATE",
      "priceCurrency": "USD",
      "unitText": "MONTH",
      "description": "Starting monthly rate per embedded engineer, minimum 3 months"
    }
  }
}
```

D. Blog post — replace the current `BlogPosting`

```
{
  "@context": "https://schema.org",
  "@graph": [
    {
      "@type": "BlogPosting",
      "@id": "https://www.brilworks.com/blog/databricks-vs-snowflake/#article",
      "headline": "Databricks vs Snowflake Compared: Features, Pricing, ML",
      "description": "REPLACE_META_DESCRIPTION",
      "image": { "@type": "ImageObject", "url": "REPLACE_BANNER_URL", "width": 1200, "height": 630 },
      "datePublished": "REPLACE_ISO_DATE",
      "dateModified": "REPLACE_ISO_DATE",
      "author": { "@id": "https://www.brilworks.com/authors/hitesh-umaletiya/#person" },
      "publisher": { "@id": "https://www.brilworks.com/#organization" },
      "mainEntityOfPage": "https://www.brilworks.com/blog/databricks-vs-snowflake/",
      "isPartOf": { "@id": "https://www.brilworks.com/#website" },
      "articleSection": "Data Engineering",
      "about": [{ "@type": "Thing", "name": "Databricks" }, { "@type": "Thing", "name": "Snowflake" }],
      "wordCount": 3200,
      "inLanguage": "en-US"
    },
```

```
    {
      "@type": "BreadcrumbList",
      "itemListElement": [
        { "@type": "ListItem", "position": 1, "name": "Home", "item": "https://www.brilworks.com/" },
        { "@type": "ListItem", "position": 2, "name": "Blog", "item": "https://www.brilworks.com/blog/" },
        { "@type": "ListItem", "position": 3, "name": "Data Engineering", "item":
"https://www.brilworks.com/blog/category/data-engineering/" },
        { "@type": "ListItem", "position": 4, "name": "Databricks vs Snowflake" }
      ]
    }
  ]
}
```

Use the real word count; the code currently multiplies reading time by 250. Also emit `og:type=article`, `article:published_time` and `article:modified_time` — all 521 posts currently have no `og:type`.

E. Author page

```
{
  "@context": "https://schema.org",
  "@type": "ProfilePage",
  "mainEntity": {
    "@type": "Person",
    "@id": "https://www.brilworks.com/authors/hitesh-umaletiya/#person",
    "name": "Hitesh Umaletiya",
    "jobTitle": "Co-founder",
    "worksFor": { "@id": "https://www.brilworks.com/#organization" },
    "url": "https://www.brilworks.com/authors/hitesh-umaletiya/",
    "image": "REPLACE_HEADSHOT_URL",
    "description": "REPLACE_WITH_2_SENTENCE_BIO_NAMING_SPECIFIC_EXPERTISE",
    "knowsAbout": ["Data engineering", "Snowflake", "Databricks", "AI agents"],
    "sameAs": ["REPLACE_LINKEDIN_URL", "REPLACE_GITHUB_OR_X_URL"]
  }
}
```

F. Case study

```
{
  "@context": "https://schema.org",
  "@type": "Article",
  "@id": "https://www.brilworks.com/case-studies/fintech-pipeline-optimization/#casestudy",
  "headline": "Cutting a fintech's pipeline costs 60% while making it 4x faster",
  "articleSection": "Case Study",
  "about": [
    { "@type": "Thing", "name": "Data pipeline optimisation" },
    { "@type": "Thing", "name": "Snowflake" },
    { "@type": "Thing", "name": "dbt" },
    { "@type": "Thing", "name": "Apache Airflow" }
  ],
  "mentions": { "@id": "https://www.brilworks.com/data-pipeline-development/#service" },
  "audience": { "@type": "BusinessAudience", "audienceType": "Fintech" },
  "abstract": "Brilworks re-architected a payments-analytics pipeline: warehouse cost down 60%, nightly run
time from 6–8 hours to about 1.5 hours.",
  "author": { "@id": "https://www.brilworks.com/#organization" },
  "publisher": { "@id": "https://www.brilworks.com/#organization" },
  "datePublished": "REPLACE_ISO_DATE",
  "dateModified": "REPLACE_ISO_DATE",
  "mainEntityOfPage": "https://www.brilworks.com/case-studies/fintech-pipeline-optimization/"
}
```

G. Free tool — `SoftwareApplication`

```
{
  "@context": "https://schema.org",
  "@type": ["SoftwareApplication", "WebApplication"],
  "name": "Snowflake Cost Calculator",
  "url": "https://www.brilworks.com/data-engineering-tools/snowflake-cost-calculator/",
  "applicationCategory": "BusinessApplication",
  "operatingSystem": "Web",
  "browserRequirements": "Requires JavaScript",
  "description": "Estimate monthly Snowflake spend from warehouse size, hours and storage.",
  "offers": { "@type": "Offer", "price": "0", "priceCurrency": "USD" },
  "publisher": { "@id": "https://www.brilworks.com/#organization" }
}
```

H. FAQPage — only where the same questions and answers are visible on the page

```
{
  "@context": "https://schema.org",
  "@type": "FAQPage",
  "mainEntity": [
    {
      "@type": "Question",
      "name": "How much do data engineering services cost?",
      "acceptedAnswer": { "@type": "Answer", "text": "REPLACE_WITH_THE_EXACT_VISIBLE_ANSWER_TEXT" }
    }
  ]
}
```

On-page meta defects found by the crawl

Defect	Count	Where
Title over 60 characters	115	48 posts, 29 tools, 8 portfolio, all 8 data tools, all 4 case studies
Meta description over 160 characters	186	98 posts, 27 tools, 11 hire pages
Two H1s	38	36 tool pages, terms page, 1 other
No H1	2	`/contact-us/`, one ebook
No `og:image`	162	All 41 tools, 31 AI agents, 17 portfolio, 17 services, 8 ebooks, 8 data tools, 4 case studies
No `og:type`	521	Every blog post
`og:type=article` on non-articles	127 pages in code	Hire, industry, tools
`og:image` URL with double slash	3	Hire pages
Weak titles	—	"About Us - Brilworks", "Staff Augmentation | Brilworks", "Brilworks | Saas Application Development Services", "Hire Dedicated Software Development Team" (no brand)
Images without alt text	108	52 in posts, 18 industry, 17 portfolio, 17 services

Part 2 — Content Architecture, E-E-A-T & Conversion

2.1 Commercial versus informational intent

What the service pages do well

- The newer Data & AI pages lead with outcomes and numbers: "cut warehouse spend 20 to 40%", "live in 7 days", "fixed-price audit".
- Most service pages run 1,100–1,450 words with five to eight question-style FAQs.
- Engagement models and pricing anchors exist (`/pricing.md` , AI agent page).

Where they miss high-intent demand

Gap	Evidence	Fix
H1s are slogans, not queries	"From raw data to running agents.", "Launch iOS & Android Apps Users Love", "Engineer enterprise data & AI on Snowflake."	Keep the slogan as a sub-headline. Put the query in the H1: "Snowflake Consulting Partner", "Mobile App Development Services"
No location or model modifiers	No page targets "offshore", "nearshore", "India", "US time-zone", "dedicated team for [industry]"	Add an "Engagement & location" block on each hire page, and build the matrix below
No industry × service pages for legacy services	Healthcare page targets "healthcare software development" only. Nothing targets "hire healthcare software developers" or "HIPAA-compliant app developers"	Build 6–8 intersection pages
Industry × data pages are thin and unlinked	`/data-engineering-for-healthtech/` is 507 words, 3 H2s, no FAQ, no canonical, no schema, 1 inbound link	Bring to service-page standard
No pricing page	Pricing exists only as `/pricing.md` for bots. Humans and Google see no pricing URL	Publish `/pricing/` with the same ranges in a table
No services hub	`/services/` is in the sitemap and linked from content but returns 404	Build it as the top of the silo tree
No comparison tables on service pages	0 of 42 service pages contain a `<table>`	Add "engagement model" and "us vs in-house vs freelancer" tables
Partner pages lack proof links	Partner pages claim certification but do not link to the official directory listing or case studies	Link out to the listing; link in from every data page
Hire pages make inconsistent promises	ReactJS page: "72 hours" in the title, "48 hours" in the H2, "7 days" on staff augmentation	Pick one verifiable SLA per page

High-intent page matrix to build (BoFu)

Query pattern	Example target	Template
hire + offshore + role	"hire offshore data engineers", "hire offshore healthcare software developers"	Hire page with location block
role + for + industry	"data engineers for fintech", "AI agent developers for healthcare"	Industry × hire intersection

Query pattern	Example target	Template
platform + consulting + partner	"Databricks consulting partner", "Snowflake migration services"	Existing partner pages, re-headed
X + migration + services	"Teradata to Snowflake migration", "Hadoop to Databricks migration"	New migration pages; blog already has the guides
service + cost	"data engineering services cost", "AI agent development cost"	`/pricing/` plus calculators
brand + vs / alternatives	"Toptal alternatives for data engineering"	Extend `/compare/` (currently 4 pages, none in sitemap)
staff aug + location	"IT staff augmentation company India"	`/staff-augmentation/` re-optimised

Cap programmatic intersections at what can carry unique proof. Each page needs at least one real project, one named technology decision and one number that appears on no other page.

2.2 Case studies and portfolio

Current state

Set	Pages	Strength	Weakness
`/case-studies/*`	4	Excellent narrative; hard numbers (60% cost, 4× speed, $18K/month bill)	740 words, no schema beyond a partial `Article`, no client attribution, 3 of 4 orphaned, hub not in sitemap, hub has no H2s
`/portfolio/*`	17 + 1 unlisted	Tech and category shown	~500 words, no schema, no `og:image`, generic titles, no outcome metrics, results gated behind a PDF form
`/x/product-listing/*`	22	—	Unclear purpose, odd URL, 1 broken, 2 wrong canonicals
`/internal/casestudies/`	1+	—	Publicly routable; should be blocked

Two parallel systems (`/case-studies/` and `/portfolio/`) split the proof.

Recommended structure

1. **One collection at `/case-studies/`.** Migrate portfolio items in and 301 the old URLs.
2. **Structured fields in the CMS**, each rendered visibly and in schema:

Field	Example	Used for
Industry	Fintech	Filter, industry-page feed, breadcrumb
Service	Data pipeline development	Service-page feed, internal link
Tech stack	Snowflake, dbt, Airflow	Tag pages, `about` in schema
Engagement model	Fixed-scope, 14 weeks, 3 engineers	Buyer qualification
Client descriptor	"Series B payments analytics platform, US"	Credibility when NDA applies
3 headline metrics	−60% cost, 4× faster, 1.5 h runtime	Hero, listing card, `abstract`

Field	Example	Used for
Before / after table	Run time, cost, incidents	LLM-extractable evidence
Client quote + role	Named where permitted	E-E-A-T
Lead engineer	Linked author profile	E-E-A-T

3. **Title formula:** `[Result] for a [client type] | [Stack] Case Study` — the fintech page already does this well; apply it to all.
4. **A summary block at the top**, in plain sentences: client, problem, what was built, result. This is the paragraph AI systems quote.
5. **Ungate the results.** Keep the PDF as an optional download, but put the numbers in HTML.
6. **Filterable hub with crawlable facets:** `/case-studies/industry/fintech/`, `/case-studies/technology/snowflake/`. Each facet page gets a 100-word intro and is self-canonical.
7. **Reciprocal links:** each service, industry and hire page shows two matching case studies; each case study ends with the matching service CTA.
8. **Target volume:** four data case studies is too few to support a Databricks and Snowflake partner claim. Aim for ten, including one per platform per priority industry.

2.3 Keyword cannibalisation and gaps

These clusters come from titles, slugs and H1s, cross-checked against GSC page data where available. The export has no query-by-page view, so confirm each cluster in Search Console (filter by query, open the Pages tab) before redirecting.

Cluster	Competing URLs	Recommended resolution
AI agent development	`/ai-agent-development-services/` and `/services/ai-agent-development/` (both in sitemap; GSC shows the duplicate at position 15 with 2,786 impressions) plus `/blog/how-to-choose-the-right-ai-agent-development-company/` and the "complete guide" post	301 `/services/ai-agent-development/` to the main service page. Posts link up with exact-match anchors
Hire Node.js developers	`/hire-nodejs-developer/` (position 68, 20,783 impressions, 9 clicks) vs 4 posts ("Node.js Hiring Guide", "Hire Node.js Developers in 2025", "Hiring Models…", "Hiring Mistakes…")	Merge the four posts into one guide; 301 the rest; link the guide to the hire page
Hire Java developers	`/hire-java-developer/` vs 3 posts	Merge to one guide
Bubble.io scaling issues	`/blog/buuble-io-scaling-issues/` (typo slug) and `/blog/bubbleio-scaling-issues/`	Merge into one, 301 the typo slug
Agentic AI in healthcare	`/blog/agentic-ai-in-healthcare/` and `/blog/agentic-ai-healthcare/`	Merge; 301
MVP development services	`/mvp-in-48-hours/` vs `/blog/mvp-development-services/`, `/blog/mvp-app-development-services/`, `/blog/what-is-mvp-development/`	Create a real `/mvp-development-services/` page; point the posts at it
Generative AI company / consulting	`/generative-ai-development-services/` vs posts titled "Generative AI Development Company" and "Generative AI Consulting Services"	Retitle the posts as buyer guides ("How to choose…") and link to the service page
Data modernization	`/data-modernization/` vs `/blog/data-modernization-services/`	Retitle the post to a how-to; link up

Cluster	Competing URLs	Recommended resolution
Cloud data engineering	`/cloud-data-engineering/` vs `/blog/cloud-data-engineering/`	Post becomes the guide; service page gets the commercial title
React vs Next.js / frameworks	`react-vs-angular`, `react-vs-vue`, `javascript-web-frameworks-comparison`, `node-js-vs-next-js`, `react-vs-nodejs`	Make the frameworks comparison the hub; each versus post links to it and to `/reactjs-development-services/`
React Native vs Flutter / cross-platform	8 posts incl. `react-native-vs-flutter` (title still says 2025), `native-vs-cross-platform`, `cross-platform-app-development-best-frameworks`, `guide-to-cross-platform-…`, `choosing-mobile-development-framework`	Keep 3 (versus post, native-vs-cross, best frameworks). Merge the rest. Refresh the year
ReactJS services vs hire ReactJS	`/reactjs-development-services/` (2 contextual links) vs `/hire-reactjs-developer/` (sitewide)	Distinct intent is fine. Differentiate titles: "React.js Development Company" vs "Hire React.js Developers"; cross-link
BI services	`/business-intelligence-services/` and `/business-intelligence-development/`	Near-duplicate intent. Keep one, 301 the other
Staffing	`/staff-augmentation/`, `/hire-dedicated-software-development-team/`, `/hire-engineers/`, `/white-label-mobile-app-development/` vs `/hire-white-label-developer/`	Define one page per intent: augmentation (individuals), dedicated team (pods), hire hub (directory)

Content gaps relative to the stated positioning

Topic	Posts today	Comment
Java	50	Largest cluster; not in the current positioning
AWS	36	Supporting
Generative AI	25	On strategy
Healthcare	21	On strategy, 13 newer compliance-focused posts are strong
AI agents	16	On strategy
Data engineering	11	**Core pillar, under-served**
Databricks	8	**Core pillar, under-served**
Snowflake	5	**Core pillar, under-served**
dbt, Airflow, data governance, data quality, MLOps	~0–3 each	**Missing**
DevOps / Kubernetes	0	Service pages exist with no supporting content

Two-thirds of the blog supports services the company now calls "legacy". The data pillar has 24 posts against Java's 50. Phase 3 addresses this.

Freshness: 40 titles or slugs reference 2025 and 8 reference 2023–2024. Several were modified in 2026 but kept the old year in the title, which lowers click-through.

Part 3 — GEO and AI Search Readiness

3.1 Entity optimisation and brand consensus

What third-party sources say (search-result level)

Source	Found	What it says	Problem
Clutch	Yes	5.0 rating; Ahmedabad; 50–249 staff; founded 2019; mobile, web, staff augmentation	Review count shown as 19 in one place and 22 in another. Service focus is legacy
GoodFirms	Yes	"Custom software developer for web, mobile and cloud… 120+ products"	Legacy positioning
Crunchbase	Yes	"IT solutions provider with a young team of experts"	Generic and outdated
G2	Yes	Unclaimed, marked inactive for over a year	Looks abandoned
Upwork, TechBehemoths, Cutshort, UpCity, SelectedFirms	Yes	Founded 2018 on some, 2019 on others; size 20–100 to 51–200	Conflicting facts
dotCMS, Storyblok partner pages	Yes	"India's leading digital transformation partner"	Off-strategy
Wikidata	Not found	—	No knowledge-graph anchor
GitHub organisation	Not found	—	No engineering footprint for an engineering firm
Databricks / Snowflake / AWS partner directories	Not found in search	Only Brilworks' own pages assert the partnership	The central claim has no independent corroboration that a model could retrieve. Verify manually

Legal-name variants in circulation: *Brilworks*, *Brilworks Software*, *Brilworks Software LLP*, *Brilworks Technology*, *Brilworks Solutions FZCO*.

The consensus problem

Brilworks' `llms.txt` tells models not to call it a "mobile app development company". But nearly every independent source, 133 of its own `og:site_name` tags, its `WebSite` schema, its `/pricing.md` ("custom software development and AI solutions company"), its About page ("Your Tech Partner… AWS Consulting Partner") and two-thirds of its blog say exactly that. Models weigh independent agreement above a self-published instruction file. Today the consensus answer to "what is Brilworks?" is "an Indian mobile and web development agency".

Missing authority signals

Signal	Action
Consistent one-line description	Use the `llms.txt` "preferred short summary" verbatim on Clutch, GoodFirms, Crunchbase, LinkedIn, G2, Upwork, partner pages, About page, schema, `og:site_name`
One founding year, one headcount range, one phone number	Decide the facts once; correct every profile
Official partner listings	Obtain the public directory URLs on Databricks, Snowflake and AWS Partner Finder; link them from the partner pages and `sameAs`
Wikidata item	Create once two or three independent, reliable sources exist (press coverage, partner directory, registry entry). Do not create it with only self-published references
GitHub organisation	Publish 3–5 small, useful repos: dbt macros, a Snowflake cost-query pack, an Airflow DAG template, the logic behind the public calculators
Founder entities	Full author pages with bio, credentials, talks, LinkedIn; 488 of 521 posts are attributed to two people whose author pages have 150 words and no schema
Named client proof	Even two named, linkable clients change how a model treats the case studies
Reviews on the data practice	The Clutch reviews describe Java, React and mobile work. Ask recent data clients for reviews that mention Snowflake or Databricks by name
Third-party mentions	Guest posts, podcast appearances and "best Databricks partners" list inclusions — the listicles are what AI answers cite for "best X company" prompts

3.2 Conversational formatting and LLM extractability

Measured across the 521 posts

Signal	Result	Assessment
Median length	2,486 words	Good depth
Has at least one question-style H2	409 (78%)	Good
Has an FAQ section	435 (83%)	Good
Contains a table	185 (35%)	Needs work — tables are the most-cited format for comparisons
"Conclusion" / "Final thoughts" style H2	340+	Generic; wasted heading
"Introduction" H2	22	Generic
Server-rendered and readable without JS	Yes	Good
Author schema connected to an entity	4	Poor
Visible "last updated" date	Yes	Good
"Summarize with AI" buttons	Yes	Nice touch; does not affect citation

Measured across service and proof pages

Signal	Result
Service pages containing a table	0 of 42
Service pages opening with a definitional sentence	Few — heroes lead with slogans
Case studies with a before/after table	0
Pricing in crawlable HTML	Partial (AI agent page); otherwise only `/pricing.md`
`/ai-agents/*` demo pages readable without JS	No — 31 pages, 2 words each
Blog index readable without JS	No

What is already good

- `llms.txt` is thorough, current (July 2026) and well structured. A second file, `llm.txt`, duplicates it in a different format; keep one.
- `robots.txt` explicitly allows GPTBot, ChatGPT-User, PerplexityBot, ClaudeBot and Google-Extended.
- Service FAQs are phrased as real buyer questions.

Formatting standard to adopt

1. **Answer-first under every H2.** The first one or two sentences answer the heading in 40–60 words, with no preamble. Detail follows.
2. **A definition sentence on every service page**, directly under the H1: "Data engineering services are…" then "Brilworks provides…". Models lift definitional sentences that name the subject.
3. **A TL;DR block** of three to five bullet points at the top of every post and case study.
4. **One table per comparison.** Any "X vs Y" or "best N" post needs a comparison table within the first screen of content.
5. **Numbers with a source and a date.** "Snowflake credits cost $2–$4 (Snowflake pricing, Oct 2026)" is quotable. "Costs vary" is not.
6. **First-party data.** Publish numbers only Brilworks has: median migration duration across its projects, average cost reduction, common failure modes. Unique statistics are the strongest citation magnet.
7. **Replace "Conclusion" headings** with a specific takeaway ("When to choose Databricks over Snowflake").
8. **Self-contained passages.** Avoid "as mentioned above"; each section should make sense when quoted alone.
9. **Server-render what you want cited.** Give each `/ai-agents/*` page a static description, sample input and output, and use-case text, or remove it from the sitemap.

3.3 Citation capture strategy

#	Action	Why it works	Owner
1	Fix entity consistency (3.1) before anything else	Models will not cite a source they cannot classify	Marketing
2	Publish 6 "reference" pages on topics Brilworks can own: Snowflake vs Databricks decision guide, Snowflake cost benchmark, Databricks migration cost model, AI agent cost model, offshore data engineering rates, embedded-team engagement models	Reference pages with tables and dated numbers are what answer engines quote	Content + engineering

#	Action	Why it works	Owner
3	Turn the 15 calculators into citable assets: publish the formula and assumptions as HTML under each tool	A calculator behind JavaScript is invisible; a published methodology is quotable	Engineering
4	Publish an annual first-party report ("State of Data Platform Costs 2027") from anonymised client data	Original data earns links and citations for a full year	Leadership
5	Add "How we work" facts pages: rates, team size, time zones, security posture, certifications, tooling	These answer the exact questions buyers put to AI assistants about vendors	Marketing
6	Get listed in third-party "best Databricks / Snowflake consulting partners" articles and directories	For "best X" prompts, assistants cite list articles, not vendor sites	Partnerships
7	Seed expert presence where models retrieve from: Stack Overflow, GitHub, Reddit data-engineering communities, dbt and Databricks community forums — as named engineers	Named expertise connects author entities to topics	Engineers
8	Build the `/compare/` section out to 10–12 pages and add it to the sitemap	Comparison prompts are the most common vendor-selection query type	Content
9	Track a fixed prompt set monthly (30 prompts across ChatGPT, Gemini, Perplexity, Claude, AI Overviews) and log mention and citation	Without a baseline, progress cannot be measured	SEO
10	Keep `llms.txt`, add `/pricing/` in HTML, and expose a clean Markdown version of key pages	Lowers extraction cost for agents	Engineering

Starter prompt set to track: "best Databricks consulting partners for mid-market", "Snowflake migration consultants", "companies that build AI agents on Snowflake data", "hire offshore data engineers", "Toptal alternatives for data engineering", "how much does a Databricks migration cost", "data engineering firms in India for US companies".

Part 4 — Prioritised Execution Roadmap

Effort: **S** under half a day · **M** 1–3 days · **L** a week or more.

Phase 1 — Quick wins (weeks 1–2)

#	Action	File / place	Effort	Impact
1	Move blog and ebook routes to ISR; return real 404s	`src/app/blog/[slug]/page.js`, `src/app/ebooks/[slug]/page.js`	M	Critical
2	Add the 12 Data & AI service pages, 4 compare pages, case-study hub and other missing pages to the sitemap	`src/app/components/lib/constants.js`	S	Critical
3	Remove from sitemap: the 404 URL, 3 `noindex` URLs, `/services/` until built	same	S	High
4	Fix 4 wrong canonicals; add canonicals to 23 pages	2 tool pages, 2 product pages, use-case and ai-solutions templates, 7 service pages	S	High
5	Link `/databricks-partner/` and `/snowflake-partner/` from the header, homepage certification section and every data page	`constants.js`, `EnterpriseCertification.js`	S	High
6	Fix footer `//` link; fix 10 dead internal links and 34 missing trailing slashes; remove `dev.` and ticket links	`Footer/index.js:198`, Storyblok content	M	High
7	Replace global schema with the linked `@graph`; fix email, phone, `sameAs`, description; remove `SearchAction`; move `LocalBusiness` off inner pages	`schemaCode.js`, `layout.js`	M	High
8	Unify `og:site_name` to "Brilworks" across 217 pages	all `page.js` metadata	S	High
9	Add `Service` schema to 40 service pages and all hire and industry pages	`generateServiceSchema` already exists	M	Medium
10	Blog meta: `og:type=article`, published/modified times, `BreadcrumbList`, real word count, author `@id`	blog template	S	Medium
11	Rewrite titles over 60 characters and descriptions over 160 on service, hire, case-study and tool pages; fix the four weak titles	metadata	M	Medium
12	One H1 on 36 tool pages; add H1 to `/contact-us/`	tool template	S	Medium
13	Add default `og:image` for tools, portfolio, case studies, ebooks	metadata	S	Low
14	Deploy the new `robots.txt`	`src/app/robots.txt/route.js`	S	Medium
15	Fix `video-sitemap.xml` 404 entries	`src/app/video-sitemap.xml/route.js`	S	Low
16	301 `/blog/author/*` and `/teams/*` to `/authors/*`; 301 `/services/ai-agent-development/` to the main page	`next.config.js`	S	Medium
17	Rewrite titles and descriptions on the high-impression, low-CTR pages in section 1.1	Storyblok	M	High
18	301 the two dead posts still earning impressions (section 1.1); `noindex` `/login/`	`next.config.js`	S	Medium

#	Action	File / place	Effort	Impact
19	Find and remove the source of the "Review snippet" appearance if it is self-serving markup	Storyblok / subdomains	S	Medium
20	Export GSC Pages (indexing) and Crawl stats; validate fixes against the reasons in section 1.1	Search Console	S	High

Phase 2 — Structural and architectural (weeks 3–6)

#	Action	Effort	Impact
1	Server-render the blog index with paginated URLs and linked category pages	L	Critical
2	Generate the sitemap from routes and CMS; split into the five-file index; add the CI guard	M	High
3	Implement the five-silo linking model: `primary_service` field on posts, related-services and related-case-studies components on every template	L	Critical
4	Backfill `primary_service` on the 293 posts with no commercial link, starting with the highest-impression posts in section 1.1	M	High
5	Build `/services/` hub, `/industries/` hub, `/pricing/` page	M	High
6	Visible, linked breadcrumbs on all templates	M	Medium
7	Performance: hero image priority on hire pages; `prefetch={false}` on nav and footer; enable image optimisation; trim third-party tags; reserve space on blog template	L	High
8	Merge `/portfolio/` into `/case-studies/` with structured fields and schema; resolve `/x/product-listing/`	L	High
9	Consolidate the confirmed cannibalisation clusters with 301s (section 2.3)	M	Medium
10	Decide the fate of `/ai-agents/*` (31), `/gist/*` (8), `/labs/` , `/posters/` : server-render real content or `noindex` and drop from sitemap	M	Medium
11	Rebalance the footer toward the five hubs and partner pages	S	Medium
12	Re-head service pages: query in the H1, definition sentence, one comparison table each	M	High
13	Rebuild author pages: bio, credentials, `ProfilePage` schema, post list	M	Medium

Phase 3 — Authority, GEO and content scale (months 2–3)

#	Action	Effort	Impact
1	Entity clean-up across Clutch, GoodFirms, Crunchbase, G2, LinkedIn, Upwork and partner pages: one description, one founding year, one headcount	M	Critical
2	Secure and link the official Databricks, Snowflake and AWS partner directory listings	M	Critical
3	Publish the six reference pages and the calculator methodologies	L	High
4	BoFu expansion led by the commercial queries in section 1.1 (AWS consulting, hire Java / Node / React, Java outsourcing) plus the data pillar: hire × industry, hire × offshore, platform migration pages, 8 more compare pages	L	High
5	Data-pillar content: 30 posts on Snowflake, Databricks, dbt, Airflow, governance, data quality, MLOps	L	High
6	Six new data case studies with named or well-described clients	L	High

#	Action	Effort	Impact
7	Refresh or merge the 48 year-stamped posts; apply the answer-first format to the top 100 posts by impressions	L	Medium
8	Launch the GitHub organisation with 3–5 repos	M	Medium
9	Request data-practice reviews on Clutch and G2; claim the G2 profile	S	Medium
10	Digital PR for inclusion in third-party partner lists; one first-party data report	L	High
11	Create the Wikidata item once independent sources exist	S	Medium
12	Monthly AI-visibility tracking against the fixed prompt set	S	Medium

How to measure success

Metric	Baseline today	90-day target
Blog posts served from cache	0 of 521	521
Dead `/blog/*` URLs returning 404	0%	100%
Indexable pages missing from sitemap	~45	0
Sitemap URLs with zero internal links	127	0
Posts with a contextual link to a service page	44%	100%
Service pages with `Service` schema	2 of 42 listed	all
Pages with wrong or missing canonical	27	0
Internal links that 404 or redirect	45 targets	0
GSC: Soft 404 + "Page with redirect" + "Duplicate without canonical"	Unknown — export the Pages report	Down 80%
GSC: clicks per day, all pages	26 (Oct 2026)	45
GSC: clicks to service, hire and industry pages	~43 per month	150 per month
GSC: US share of clicks	7%	15%
GSC: average position of the 19 hire pages	45–68	under 30
GSC: clicks from Snowflake / Databricks / data-engineering queries	7 in 16 months	100 per month
AI prompt set: Brilworks mentioned	Unknown — run baseline	Set after baseline

Appendix A — Additional observations

- **Security:** `src/app/components/Teams/TeamAuthor.js` contains a hard-coded Storyblok token used with `version=draft` in browser code. If that is a preview token, unpublished content is readable by anyone. Rotate it and move the call to the server.
- **Publicly routable internal pages:** `/internal/casestudies/`, `/admin/consent-logs/`, `/dashboard/` resolve on the public site. `admin` is `noindex`; all should be disallowed and access-controlled.

- **Security headers:** HSTS is present. No Content-Security-Policy, X-Frame-Options, X-Content-Type-Options or Referrer-Policy on HTML responses.
- **Homepage cache age:** the prerendered homepage was 20 days old at crawl time, so the two "latest posts" it shows are stale.
- **About page:** contains the stat "60% Professional Employee", which appears to be a typo, and 535 words with no founder names, history or credentials in the server HTML.

Appendix B — Method

- Crawl: all 726 URLs in `sitemap.xml`, plus 153 URLs that were linked internally or present in the codebase but not in the sitemap. Collected status, redirect chain, cache headers, title, description, canonical, robots, headings, JSON-LD, images, word counts and links.
- Link graph: built from server HTML only, which is what non-rendering crawlers see.
- Lab performance: Chromium with mobile emulation, 4× CPU throttling, slow-4G network profile. Third-party hosts and CMS image hosts were not reachable from the audit environment.
- Entity research: web search results only.
- Search Console: Performance export supplied by Brilworks (Web search, 10 Jul 2025 – 9 Nov 2026). Query and page files are capped at 1,000 rows by Google.

Appendix C — Sources for third-party entity findings

- Clutch: https://clutch.co/profile/brilworks-software/
- GoodFirms: https://www.goodfirms.co/company/brilworks-software-llp
- Crunchbase: https://www.crunchbase.com/organization/brilworks
- G2: https://www.g2.com/products/brilworks-software/pricing
- Upwork: https://www.upwork.com/ag/brilworks/
- TechBehemoths: https://techbehemoths.com/company/brilworks-software-llp
- UpCity: https://upcity.com/profiles/brilworks-software
- dotCMS partner page: https://auth.dotcms.com/partners/brilworks
- Storyblok partner page: https://storyblok.com/pl/brilworks-technology
- YourStory: https://yourstory.com/companies/brilworks/amp
- TechReviewer: https://techreviewer.co/companies/brilworks-software